

Logistic_Regression

June 19, 2021

```
[19]: import sys
      !{sys.executable} -m pip install ipynb

      import warnings
      warnings.filterwarnings('ignore')
```

Requirement already satisfied: ipynb in
c:\users\gampl\appdata\local\programs\python\python39\lib\site-packages (0.5.1)

WARNING: You are using pip version 21.1.1; however, version 21.1.2 is available.
You should consider upgrading via the
'c:\users\gampl\appdata\local\programs\python\python39\python.exe -m pip install
--upgrade pip' command.

```
[20]: %%capture

      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      from ipynb.fs.full.CleaningData import getDataset
      from ipynb.fs.full.CleaningData import getCovarianceVector
      from sklearn.model_selection import train_test_split
      from copy import deepcopy

      pd.set_option("display.max_rows", 10000)
      pd.set_option("display.max_columns", None)
```

```
[ ]: %%time
      df = getDataset(500)
```

```
[ ]: %%time
      covVec = getCovarianceVector(df)

      fig = plt.figure(figsize=(16, 6)) # the figsize changes the width and height,
      ↪ respectively
      ax = fig.add_axes([0,0,1,1])
      langs = covVec.index
      students = covVec
      ax.bar(langs,students, align='edge', width=0.7) #width determines width of bars
```

```
plt.xticks(rotation = 90)
# plt.show()
```

```
[ ]: # Separating and storing y values (Status Feature)
```

```
df.head(10)

statusColumn = df['status']
del df['status']
urlColumn = df['url']
del df['url']

# Printing the feature set

print(df.columns)

# Storing the number of features

numFeatures = (len(df.columns))
# print(numFeatures)
```

```
[24]: theta = np.random.uniform(-0.5,0.5,numFeatures) # Initialising random theta
      ↪ values
origTheta = theta # Keeping a copy of original theta values

# Splitting data into test, validation and train
trainX, testX, trainY, testY = train_test_split(df, statusColumn, test_size=0.
      ↪ 15)
trainX, validationX, trainY, validationY = train_test_split(trainX, trainY,
      ↪ test_size=0.2)
```

```
[25]: %%time
def sigmoid(x,k): #Sigmoid function

    if(k==0):
        f = 1
    else:
        f = k
    return 1/(1+np.exp(-f*x))

def prediction(theta, x, k): #IPrediction function
    return sigmoid(np.matmul(np.transpose(theta),x),k)

def LogisticRegression(trainX, trainY, theta, a, regularisation): # Logistic
      ↪ Regression Algorithm
    alpha = a
    newTheta = theta
```

```

strengthOfRegularisation = regularisation
count = 0

while(count < 10):
    oldTheta = newTheta
    for i in trainX.index:
        for j in range(numFeatures):
            newTheta[j] = newTheta[j] + alpha*(trainY.loc[i] -
↪prediction(newTheta,trainX.loc[i],strengthOfRegularisation))*(trainX.
↪loc[i][j]) + alpha*strengthOfRegularisation*newTheta[j]
            count = count + 1

    return newTheta

def calculateAccuracy(x,y,theta,confusionMatrix, k): # Calculates the accuracy
↪and create the confusion matrix
    numberCorrect = 0
    numberIncorrect = 0

    for i in range(len(testY)):

        predicted_result = round(prediction(newTheta,testX.iloc[i], k))

        if(predicted_result == 1 and testY.iloc[i] == 0):
            confusionMatrix[0] = confusionMatrix[0] + 1
        else:
            if (predicted_result == 0 and testY.iloc[i] == 1):
                confusionMatrix[1] = confusionMatrix[1] + 1
            else:
                if(predicted_result == 1 and testY.iloc[i] == 1):
                    confusionMatrix[2] = confusionMatrix[2] + 1
                else:
                    if(predicted_result == 0 and testY.iloc[i] == 0):
                        confusionMatrix[3] = confusionMatrix[3] + 1

        if(predicted_result == testY.iloc[i]):
            numberCorrect = numberCorrect + 1
        else:
            numberIncorrect = numberIncorrect + 1

    averageCorrect = numberCorrect/len(testY)
    averageIncorrect = numberIncorrect/len(testY)

    return averageCorrect*100, averageIncorrect*100, confusionMatrix

```

Wall time: 0 ns

```

[26]: %%time

results = [] # Matrix to store validation results
confusionMatrix = [0,0,0,0] # Confusion Matrix
alphaArray = [0.01, 0.001, 0.0001] # Alpha Values
regularisationArray = [0, 0.001, 0.002, 0.01, 0.03] # Regularisation values
newTheta = []

theta = np.random.uniform(-0.5,0.5,numFeatures) # Initialising random theta
↪values
origTheta = deepcopy(theta) # Keeping a copy of original theta values

# Obtaining the optimal hyperparameters by getting the accuracy and confusion
↪matrix for each hyperparameter

r = 0
for i in range(3):
    for j in range(5):
        newTheta = LogisticRegression(validationX, validationY, theta,
↪alphaArray[i], regularisationArray[j])
        results.append(calculateAccuracy(testX,testY,newTheta,confusionMatrix,
↪0))

        print("Alpha = ",alphaArray[i])
        print("Regularisation = ", regularisationArray[j])
        print("=====")
        print("Confusion Matrix",'\n' )
        print(results[r][2][2],"    ",results[r][2][1] )
        print(results[r][2][0],"    ",results[r][2][3] )
        print("=====")
        print("Percentage Correct : ", results[r][0])
        print("Percentage Incorrect : ", results[r][1])
        print("=====",'\n')
        theta = deepcopy(origTheta)
        r = r + 1
        confusionMatrix = [0,0,0,0]

```

```

Alpha = 0.01
Regularisation = 0
=====
Confusion Matrix

815    78
33     797
=====
Percentage Correct : 93.55774811375508
Percentage Incorrect : 6.4422518862449225
=====

```

```

Alpha = 0.01
Regularisation = 0.001
=====
Confusion Matrix

739      154
57       773
=====
Percentage Correct : 87.7539175856065
Percentage Incorrect : 12.2460824143935
=====

```

```

Alpha = 0.01
Regularisation = 0.002
=====
Confusion Matrix

748      145
55       775
=====
Percentage Correct : 88.39233894370284
Percentage Incorrect : 11.607661056297156
=====

```

```

Alpha = 0.01
Regularisation = 0.01
=====
Confusion Matrix

784      109
51       779
=====
Percentage Correct : 90.71387115496228
Percentage Incorrect : 9.286128845037725
=====

```

```

Alpha = 0.01
Regularisation = 0.03
=====
Confusion Matrix

787      106
41       789
=====
Percentage Correct : 91.4683691236216
Percentage Incorrect : 8.53163087637841
=====

```

```

Alpha = 0.001
Regularisation = 0
=====
Confusion Matrix

806      87
31       799
=====
Percentage Correct : 93.15147997678469
Percentage Incorrect : 6.848520023215323
=====

```

```

Alpha = 0.001
Regularisation = 0.001
=====
Confusion Matrix

732      161
55       775
=====
Percentage Correct : 87.46372605919908
Percentage Incorrect : 12.53627394080093
=====

```

```

Alpha = 0.001
Regularisation = 0.002
=====
Confusion Matrix

733      160
54       776
=====
Percentage Correct : 87.57980266976205
Percentage Incorrect : 12.420197330237958
=====

```

```

Alpha = 0.001
Regularisation = 0.01
=====
Confusion Matrix

739      154
53       777
=====
Percentage Correct : 87.98607080673244
Percentage Incorrect : 12.013929193267558
=====

```

```

Alpha = 0.001
Regularisation = 0.03
=====
Confusion Matrix

748      145
51        779
=====
Percentage Correct : 88.62449216482878
Percentage Incorrect : 11.375507835171213
=====

```

```

Alpha = 0.0001
Regularisation = 0
=====
Confusion Matrix

661      232
78        752
=====
Percentage Correct : 82.0081253627394
Percentage Incorrect : 17.99187463726059
=====

```

```

Alpha = 0.0001
Regularisation = 0.001
=====
Confusion Matrix

669      224
67        763
=====
Percentage Correct : 83.11085316308764
Percentage Incorrect : 16.88914683691236
=====

```

```

Alpha = 0.0001
Regularisation = 0.002
=====
Confusion Matrix

669      224
67        763
=====
Percentage Correct : 83.11085316308764
Percentage Incorrect : 16.88914683691236
=====

```

```

Alpha = 0.0001
Regularisation = 0.01
=====
Confusion Matrix

670      223
67       763
=====
Percentage Correct : 83.16889146836913
Percentage Incorrect : 16.831108531630875
=====

```

```

Alpha = 0.0001
Regularisation = 0.03
=====
Confusion Matrix

678      215
69       761
=====
Percentage Correct : 83.51712130005804
Percentage Incorrect : 16.48287869994196
=====

```

Wall time: 51min 22s

```

[28]: %%time
trainingResults = []
confusionMatrix = [0,0,0,0]
newTheta = []

theta = np.random.uniform(-0.5,0.5,numFeatures) # Initialising random theta
↳ values
origTheta = deepcopy(theta) # Keeping a copy of original theta values

# Obtaining the optimal values for theta by getting the accuracy and confusion
↳ matrix using the tuned hyperparameters

newTheta = LogisticRegression(trainX, trainY, theta, 0.01, 0)
trainingResults.append(calculateAccuracy(trainX, trainY, newTheta,
↳ confusionMatrix, 0))

```

Wall time: 13min 19s

```

[30]: %%time
# Obtaining the test results by using new data on the model with trained theta
↳ values

```



```

confusionMatrix = [0,0,0,0]

testResults = []
testResults.append(calculateAccuracy(testX, testY, newTheta, confusionMatrix,
    ↪0))
print("Alpha = ",0.01)
print("Regularisation = ", 0)
print("=====")
print("Confusion Matrix",'\n' )
print(testResults[0][2][2],"    ",testResults[0][2][1] )
print(testResults[0][2][0],"    ",testResults[0][2][3] )
print("=====")
print("Percentage Correct : ", testResults[0][0])
print("Percentage Incorrect : ", testResults[0][1],)
print("=====",'\n')

```

```

Alpha =  0.01
Regularisation =  0
=====
Confusion Matrix

820      73
34      796
=====
Percentage Correct :  93.78990133488102
Percentage Incorrect :  6.210098665118979
=====

Wall time: 184 ms

```

[]: